

Introduction to transaction processing

An interesting experiment

```
$sqlplus jrg@csci
```

```
SQL> SELECT COUNT(*)
2 FROM SKILL;

COUNT (*)
-----
19
```

```
$sqlplus jrg@csci
```

```
SQL> SELECT COUNT(*)
2 FROM SKILL;

COUNT (*)
-----
19
```

An interesting experiment

```
SQL> INSERT INTO SKILL
2 VALUES('singing');
1 row created.
```

```
SQL> SELECT COUNT(*)
2 FROM SKILL;

COUNT (*)
-----
20
```

```
SQL> SELECT COUNT(*)
2 FROM SKILL;

COUNT (*)
-----
19
```

An interesting experiment

```
SQL> COMMIT;
Commit complete.
```

```
SQL> SELECT COUNT(*)
2 FROM SKILL;

COUNT (*)
-----
20
```

```
SQL> SELECT COUNT(*)
2 FROM SKILL;

COUNT (*)
-----
20
```

Transaction ? What is it ?

A partially ordered set of *read*, *write* operations on the database items is called as a transaction

Principles of transaction processing

Users interact with the database by executing programs

Execution of a program is equivalent to execution of a partially ordered set of *read*, *write* operations
Database is visible to transactions as a collection of data items

Concurrently running transactions interleave their operations

Transactions have no impact on execution of their operations

Principles of transaction processing

Each transaction terminates by either *commit* or *abort* operation

Each transaction arrives at a consistent database state and must leave a database in a consistent state as well

So, where is a problem ?

T ₁	T ₂	x = \$100
read(x)		x = \$100
	read(x)	x = \$100
write(x, x-10)		x = \$90
	write(x, x+20)	x = \$120
	commit	x = \$120
commit		x = \$120


Correctness condition

What makes concurrent execution of database transaction incorrect ?


How do we define a correct concurrent execution of database transactions ?

Conflicting operations

	read	write
read	NO	YES
write	YES	YES

**Conflict serializability condition**

Concurrent execution of database transactions is conflict serializable if there exists a possible serial execution of the same set of transactions such that in both executions the order of conflicting operations is the same

Conflict serializable execution

Order of conflicting operations: T₁ before T₂

0 Introduction to transaction processing

Conflict nonserializable execution

T ₁	T ₂	x = \$100
read(x)		x = \$100
	read(x)	x = \$100
write(x, x-10)		x = \$90
	write(x, x+20)	x = \$120
	commit	x = \$120
commit		x = \$120

Order of conflicting operations:
T₁ before T₂ and T₂ before T₁ **impossible to serialize**

© Janusz R. Getta CSC1235/MCS9235/CSC1835 Databases, SCSSE, Autumn 2015 13

0 Introduction to transaction processing

Serialization graph

© Janusz R. Getta CSC1235/MCS9235/CSC1835 Databases, SCSSE, Autumn 2015 14

0 Introduction to transaction processing

Serialization graph

© Janusz R. Getta CSC1235/MCS9235/CSC1835 Databases, SCSSE, Autumn 2015 15

0 Introduction to transaction processing

Serialization graph testing protocol (SGT)

Principles
Scheduler maintains and tests serialization graph
If an operation issued by a transaction violates conflict serializability (i.e. it creates a cycle in serialization graph) then such transaction is aborted

Problems
Cascading aborts
performance (testing acyclicity of serialization graph has O(n²) complexity)

© Janusz R. Getta CSC1235/MCS9235/CSC1835 Databases, SCSSE, Autumn 2015 16

0 Introduction to transaction processing

Two-phase locking (2PL) protocol

Principle
Each transaction must acquire all locks before releasing any lock

Problems
Deadlocks
Unnecessary locks when execution is conflict serializable

© Janusz R. Getta CSC1235/MCS9235/CSC1835 Databases, SCSSE, Autumn 2015 17

0 Introduction to transaction processing

Two-phase locking (2PL) protocol

T ₁	T ₂
lock(u) read(u)	
	lock(v) write(v, v+1)
write(u, u+2)	
lock(v) wait	
	lock(x) read(x)
	unlock(v)
	write(x, x+2)
...	

© Janusz R. Getta CSC1235/MCS9235/CSC1835 Databases, SCSSE, Autumn 2015 18

0 Introduction to transaction processing	
Two-phase locking (2PL) protocol	
...	
T ₁	T ₂
lock (v)	
	unlock (x)
write (v, v+1)	
unlock (v)	
unlock (u)	

0 Introduction to transaction processing	
Deadlock	
T ₁	T ₂
lock (u) read (u)	
	lock (v) write (v, v+1)
lock (v) wait	
	lock (u) wait

0 Introduction to transaction processing	
Timestamp ordering (TO) protocol	
Principles	
Each transaction obtains a timestamp at the start point	
Data items are stamped each time any transaction accesses data items in a read or write mode	
Access to data items is permitted in increasing order of timestamps	
Problems	
Cascading aborts	
Unnecessary aborts when execution is conflict serializable	

0 Introduction to transaction processing		
Timestamp ordering (TO) protocol		
T ₁	T ₂	x
timestamp(t1)		
read(x)		x:t1
write(x, x-10)		
	timestamp(t2)	
	write(x)	x:t1:t2
	read(y)	y:t2
read(y)		y:t2:t1
write(y, y+1)		y:t2:t1
abort		

© Janusz R. Getta CSC235/MCS9235/CSC835 Databases, SCSSE, Autumn 2015 22

0 Introduction to transaction processing

Timestamp ordering (TO) protocol

T_1	T_2	x
timestamp(t1)		
read(x)		x:t1
write(x, x-10)		
	timestamp(t2)	
	read(x)	x:t1:t2
fail		
	forced abort	

Cascading abort !

© Janusz R. Getta

CS2235/MCS9235/CS835 Databases, SCSSSE, Autumn 2015

23

Cascading abort !

0 Introduction to transaction processing	
References	
Elmasri R., Navathe S., Fundamentals of Database Systems, 6th edition, chapter 21 Introduction to Transaction Processing Concepts and Theory, pp. 747-779	
Elmasri R., Navathe S., Fundamentals of Database Systems, 6th edition, chapters 22.1, 22.3 Concurrency Control Techniques, pp. 780-794	