

Introduction to transaction processing

An interesting experiment

```
$sqlplus jrg@csci
```

```
SQL> SELECT COUNT(*)
2 FROM SKILL;

COUNT (*)
-----
19
```

```
$sqlplus jrg@csci
```

```
SQL> SELECT COUNT(*)
2 FROM SKILL;

COUNT (*)
-----
19
```

An interesting experiment

```
SQL> INSERT INTO SKILL
2 VALUES('singing');
1 row created.
```

```
SQL> SELECT COUNT(*)
2 FROM SKILL;

COUNT (*)
-----
20
```

```
SQL> SELECT COUNT(*)
2 FROM SKILL;

COUNT (*)
-----
19
```

An interesting experiment

```
SQL> COMMIT;
Commit complete.
```

```
SQL> SELECT COUNT(*)
2 FROM SKILL;

COUNT (*)
-----
20
```

```
SQL> SELECT COUNT(*)
2 FROM SKILL;

COUNT (*)
-----
20
```

Transaction ? What is it ?

A partially ordered set of *read*, *write* operations on the database items is called as a transaction

Principles of transaction processing

Users interact with the database by executing programs

Execution of a program is equivalent to execution of a partially ordered set of *read*, *write* operations

Database is visible to transactions as a collection of data items

Concurrently running transactions interleave their operations

Transactions have no impact on execution of their operations

Principles of transaction processing

Each transaction terminates by either *commit* or *abort* operation

Each transaction arrives at a consistent database state and must leave a database in a consistent state as well

So, where is a problem ?

| T ₁ | T ₂ | x = \$100 |
|----------------|----------------|-----------|
| read(x) | | x = \$100 |
| | read(x) | x = \$100 |
| write(x, x-10) | | x = \$90 |
| | write(x, x+20) | x = \$120 |
| | commit | x = \$120 |
| commit | | x = \$120 |


Correctness condition

What makes concurrent execution of database transaction incorrect ?


How do we define a correct concurrent execution of database transactions ?

Conflicting operations

| | read | write |
|-------|------|-------|
| read | NO | YES |
| write | YES | YES |

**Conflict serializability condition**

Concurrent execution of database transactions is conflict serializable if there exists a possible serial execution of the same set of transactions such that in both executions the order of conflicting operations is the same

Conflict serializable execution

Order of conflicting operations: T₁ before T₂

0 Introduction to transaction processing

Conflict nonserializable execution

| T ₁ | T ₂ | x = \$100 |
|----------------|----------------|-----------|
| read(x) | | x = \$100 |
| | read(x) | x = \$100 |
| write(x, x-10) | | x = \$90 |
| | write(x, x+20) | x = \$120 |
| | commit | x = \$120 |
| commit | | x = \$120 |

Order of conflicting operations:
T₁ before T₂ and T₂ before T₁ **impossible to serialize**

© Janusz R. Getta CSCI235/MCS9235/CSCI835 Databases, SCSSE, Autumn 2015 13

0 Introduction to transaction processing

Serialization graph

© Janusz R. Getta CSCI235/MCS9235/CSCI835 Databases, SCSSE, Autumn 2015 14

0 Introduction to transaction processing

Serialization graph

© Janusz R. Getta CSCI235/MCS9235/CSCI835 Databases, SCSSE, Autumn 2015 15

0 Introduction to transaction processing

Serialization graph testing protocol (SGT)

Principles
Scheduler maintains and tests serialization graph
If an operation issued by a transaction violates conflict serializability (i.e. it creates a cycle in serialization graph) then such transaction is aborted

Problems
Cascading aborts
performance (testing acyclicity of serialization graph has O(n²) complexity)

© Janusz R. Getta CSCI235/MCS9235/CSCI835 Databases, SCSSE, Autumn 2015 16

0 Introduction to transaction processing

Two-phase locking (2PL) protocol

Principle
Each transaction must acquire all locks before releasing any lock

Problems
Deadlocks
Unnecessary locks when execution is conflict serializable

© Janusz R. Getta CSCI235/MCS9235/CSCI835 Databases, SCSSE, Autumn 2015 17

0 Introduction to transaction processing

Two-phase locking (2PL) protocol

| T ₁ | T ₂ |
|---------------------|-----------------------|
| lock(u) read(u) | |
| | lock(v) write(v, v+1) |
| write(u, u+2) | |
| lock(v) wait | |
| | lock(x) read(x) |
| | unlock(v) |
| | write(x, x+2) |
| ... | |

© Janusz R. Getta CSCI235/MCS9235/CSCI835 Databases, SCSSE, Autumn 2015 18

| 0 Introduction to transaction processing | |
|--|----------------|
| Two-phase locking (2PL) protocol | |
| ... | |
| T ₁ | T ₂ |
| lock (v) | |
| | unlock (x) |
| write (v, v+1) | |
| unlock (v) | |
| unlock (u) | |

| 0 Introduction to transaction processing | |
|--|-------------------------|
| Deadlock | |
| T ₁ | T ₂ |
| lock (u) read (u) | |
| | lock (v) write (v, v+1) |
| lock (v) wait | |
| | lock (u) wait |

| 0 Introduction to transaction processing | |
|--|--|
| Timestamp ordering (TO) protocol | |
| Principles | |
| Each transaction obtains a timestamp at the start point | |
| Data items are stamped each time any transaction accesses data items in a read or write mode | |
| Access to data items is permitted in increasing order of timestamps | |
| Problems | |
| Cascading aborts | |
| Unnecessary aborts when execution is conflict serializable | |

| 0 Introduction to transaction processing | | |
|--|----------------|---------|
| Timestamp ordering (TO) protocol | | |
| T ₁ | T ₂ | x |
| timestamp (t1) | | |
| read(x) | | x:t1 |
| write(x, x-10) | | |
| | timestamp (t2) | |
| | write(x) | x:t1:t2 |
| | read(y) | y:t2 |
| read(y) | | y:t2:t1 |
| write(y, y+1) | | y:t2:t1 |
| abort | | |

© Janusz R. Getta CSC235/MCS9235/CSC8135 Databases, SCSSE, Autumn 2015 22

0 Introduction to transaction processing

Timestamp ordering (TO) protocol

| T_1 | T_2 | x |
|----------------|---------------|---------|
| timestamp(t1) | | |
| read(x) | | x:t1 |
| write(x, x-10) | | |
| | timestamp(t2) | |
| | read(x) | x:t1:t2 |
| fail | | |
| | forced abort | |

Cascading abort !

© Janusz R. Getta

CS223S/MCS923S/CS183S Databases, SCSS, Autumn 2015

23

Cascading abort !

| 0 Introduction to transaction processing | |
|---|--|
| References | |
| Elmasri R., Navathe S., Fundamentals of Database Systems, 6th edition, chapter 21 Introduction to Transaction Processing Concepts and Theory, pp. 747-779 | |
| Elmasri R., Navathe S., Fundamentals of Database Systems, 6th edition, chapters 22.1, 22.3 Concurrency Control Techniques, pp. 780-794 | |